

Improving Particle Swarm Optimization in a Parallel World

Matthew Gardner

April 11, 2009

1 Introduction

Optimization problems are ubiquitous in our modern world. Businesses such as Google need to decide where to place ads, scientists need to fit models to data, and airlines need to schedule their flights. All of those problems are optimization problems, and optimization techniques can be used to find solutions. Optimization algorithms have been derived from all sorts of natural phenomena. Some methods look at the landscape of the function as they are searching and follow hills to the best values, as in the gradient ascent method. Other methods take ideas from biology or metallurgy, like genetic algorithms and simulated annealing. Particle swarm optimization is a recently developed optimization technique that draws on ideas from the sociology of flocking birds.

All of these methods, however, are fundamentally sequential in nature—they need to be run in a specific order, and because of that, they are typically only run on one machine. The problems people want to solve are getting bigger, and larger problems need to run on multiple machines in parallel. Parallel implementations of these algorithms have been attempted, but they do not often perform as well as sequential versions for the amount of computation performed. This work focuses on one of those algorithms, particle swarm optimization, and improving its performance in a large-scale, parallel environment. The work I describe in Section 3 was done in collaboration with other students I work with. The work in Section 4 I did by myself.

The rest of this paper is outlined as follows. In Section 2, I formally describe the particle swarm optimization algorithm, after giving a brief description of optimization in general. In Section 3, I describe the sociological aspect of particle swarm optimization and improvements that we have made for large swarms. Section 4 then introduces the idea of speculative execution in general and shows how it can be done in particle swarm optimization. I then conclude in Section 5.




Figure 1: A complicated function: $\cos(x) \sin(10x) e^{-|x|}$

2 Particle Swarm Optimization

2.1 Optimization in General

The goal of optimization is to find the parameters to a function that either maximize or minimize the function value. To take an example from above, Google needs to find the placements of website advertisements that generate the most clicks for those websites. In this example, the placements of advertisements are the parameters to the function, and the number of clicks generated is the function to be optimized. A more mathematical example can be seen in Figure 1. The function shown has a point x for which $f(x)$ is the highest. Finding the precise location of that point is not a trivial task, especially for more complicated functions, and it is that problem which optimization techniques try to solve.

2.2 Particle Swarm Optimization Specifically

Particle swarm optimization was proposed in 1995 by James Kennedy and Russell Eberhart. It tries to intelligently search a multi-dimensional space by mimicking the swarming and flocking behavior of birds and other animals [4]. It is a sociological algorithm that depends on interaction between particles to quickly and consistently find the optimal solution to a problem. The algorithm keeps track of a number of potential solutions, called particles, which move somewhat randomly through the search space at each iteration. When the algorithm is done, the best position that any of the particles ever found is returned. Particles remember

the best place they have been, or solution they have evaluated, and when they move they are attracted back to that place, as well as to the best solution other particles have seen. Specifically, the formulas for updating the position \vec{x}_t and velocity \vec{v}_t of a particle at iteration t are as follows:

$$\vec{v}_{t+1} = \chi [\vec{v}_t + \phi_1 U() \otimes (\vec{p} - \vec{x}_t) + \phi_2 U() \otimes (\vec{g} - \vec{x}_t)] \quad (1)$$

$$\vec{x}_{t+1} = \vec{x}_t + \vec{v}_{t+1} \quad (2)$$

where $U()$ is a vector of random numbers drawn from a uniform distribution, the \otimes operator is an element-wise vector multiplication, \vec{p} (called pbest) is the best position the current particle has seen, and \vec{g} (called gbest) is the best position any of the other particles have seen. ϕ_1 , ϕ_2 , and χ are parameters with prescribed values required to ensure convergence (2.05, 2.05, and .73, respectively) [2] [7].

The sociology in this algorithm defines which other particles form the “neighborhood” of a given particle—the other particles whose best solution it sees (again, that is \vec{g} in the equations above). If particle 1 is a neighbor of particle 2, then particle 1 will be influenced by good values that particle 2 sees. But if particle 3 is not a neighbor of particle 1, particle 1 will be completely unaware of good values that particle 3 finds.

There are many ways to define a particle’s neighborhood, varying from the entire rest of the swarm to just one other particle.

The way a neighborhood is defined can have drastic effects on the performance of the algorithm; the more neighbors each particle has, the faster information spreads throughout the swarm, and the quicker the algorithm converges. For some problems it is possible that the algorithm converges too quickly and gets caught in a local optimum; it stops on a hill when there is a mountain next to it. For those kinds of problems, less communication, or smaller neighborhoods, is often preferable.

3 Topology in PSO

Much work has been done on the sociological aspect of particle swarm optimization, trying to determine the best neighborhood structure to use for various classes of problems. That work focused on swarms of tens of particles to several hundred particles on a single machine. Conclusions were reached about which neighborhood structures worked best, but only in this limited context [1]. With the advent of parallel processing and the potential to have swarms of thousands or hundreds of thousands of particles, some of the earlier work needs rethinking. This section focuses on neighborhood structures in very large swarms, especially when the algorithm is run in parallel on multiple machines.

3.1 Defining Topology

Many different neighborhood structures have been proposed and tested [5]. In the literature, neighborhood structures are referred to as topologies. Two of them are more common—the




Figure 2: The complete and ring topologies in graph form. Figure taken from tracer.uc3m.es/tws/pso/neighborhood.html.

complete or gbest topology, where all particles communicate with all other particles, and the ring or lbest topology, where each particle only communicates with its two immediate neighbors [1].

To formalize the description of topologies, we can consider the swarm as a graph $\{V, E\}$, where the set of particles forms the nodes V on the graph and the (directed) edges E between nodes represent communication between particles. If $(1,2)$ is an edge in the graph, then particle 1 influences particle 2's gbest, or particle 1 is in particle 2's neighborhood. The neighborhood N_i of particle i can be written formally as

$$N_i = \{p_j | (j, i) \in E\}$$

Neighborhoods can also be thought of in terms of which other particles a particular particle sends its information to, instead of which ones it gets information from. This conceptualization makes the communication a little bit more explicit (and more directly relates to how to actually run the algorithm) and will be used to describe topologies throughout the rest of this paper. N_i then becomes

$$N_i = \{p_j | (i, j) \in E\}$$

The two topologies mentioned previously, complete and ring, would then have the following neighborhoods for each particle i in a swarm of n particles:

$$\text{Complete : } N_i = \{1, 2, 3, \dots, n\}$$

$$\text{Ring : } N_i = \{i - 1, i, i + 1\}$$

For the benefit of more visual readers, the two topologies are shown in graph form in Figure 2; notice how many more lines are in the complete topology than in the ring topology.

3.2 Which Topology to Use

There is a theorem about optimization called “No Free Lunch” [8]. In essence it says that no optimization algorithm can do better than any other optimization algorithm across all functions. If you do better on some functions with a given algorithm that is because you do worse on other functions. This applies directly to the question of which topology in PSO is the “best,” because it proves that there really is no “best” topology across all functions.




Figure 3: PSO performance on the function Rastrigin with varying numbers of particles.

When selecting a topology you must have a specific type of function in mind, and pick the topology that works best for that type of function.

An interesting question along those lines is, “Are the problems that most people are interested in solving inherently similar?” If that is the case, one could design an algorithm that performs better on the problems people care about, taking a performance hit only on functions that don’t matter anyway. But, because of No Free Lunch, any argument claiming that a particular algorithm or topology is better than another must make that claim only for a specific set of functions.

When deciding which topology to use, then, making a wise decision requires having a certain kind of function in mind.

The complete topology is often preferable for many problems, because it allows information to be shared quickly throughout the swarm and helps the algorithm converge more quickly. However, too much shared information can cause premature convergence, as mentioned previously, so some researchers are wary of the complete topology. In the absence of knowledge about the function to be optimized, researchers recommended, the ring topology should be used. It might take longer, but at least it won’t get stuck, they say [1].

For a swarm of 50 particles it is accurate that complete often doesn’t perform well, and most researchers were satisfied to stop with a swarm that size. We have shown, however, that increasing the number of particles in a swarm will help the algorithm avoid premature convergence, making complete the best topology even for somewhat complicated functions [6]. Figure 3 shows the performance of PSO on a commonly used benchmark function, the Rastrigin function, with varying numbers of particles. Notice that with 50 particles the algorithm

does get stuck pretty early, not finding significantly better values after 250 iterations. This is precisely what many researchers in the field were saying. With more particles, however, the algorithm finds much better values and continues to find better values well past 250 iterations.

3.3 Communication Overhead

In a parallel system, another problem with the complete topology is that communication between processors is expensive relative to the amount of time it takes to evaluate the functions many people work with. A quick look at Figure 2 shows the incredibly large number of messages that need to be sent between particles for a relatively small swarm. Communicating all of the information from every particle at every iteration can take orders of magnitude more time than doing a simple evaluation and render parallelization of the algorithm worthless.

We developed a modification to the algorithm along with a new topology that approximates the complete topology with much less communication. If each particle at each iteration randomly chooses a small set of particles to send its information to, eventually information will be communicated throughout the whole swarm. We call this the Random topology, and it is formally defined as follows:

$$\text{Random} : N_i = \{i, U_1, U_2, \dots, U_k\}$$

where U_j is a random number from 1 to $n - 1$, and k is the number of neighbors to send information to at each iteration.

In the standard particle swarm optimization algorithm, the best neighbor that a particle is attracted to (the \vec{g} in the formula above) is taken from its current neighborhood. With the random topology we proposed, particles can get confused, being attracted to a different neighbor at every iteration, and performance significantly decreases. We modified the algorithm slightly such that each particle remembers the best neighbor it has ever had and is attracted to that position instead of its current best neighbor. Using this modification and a random topology, we saw performance that was indistinguishable from complete with only 5% of the communication, and the speed up is significant even in a serial implementation [6].

Figure 4 shows the performance of Random versus Complete on the function Rastrigin for various swarm sizes. The x-axis shows the number of particles in the swarm, and the y-axis shows the best function value found after 500 iterations. It can be seen from the graph that Random does a very good job at approximating the Complete topology while required only a small fraction of the communication.

3.4 Subswarms

We have run all of the experiments I've talked about so far with swarms of up to 4000 particles. All of them were run with a serial implementation of particle swarm optimization on a single machine, though the results can be used in designing parallel algorithms.

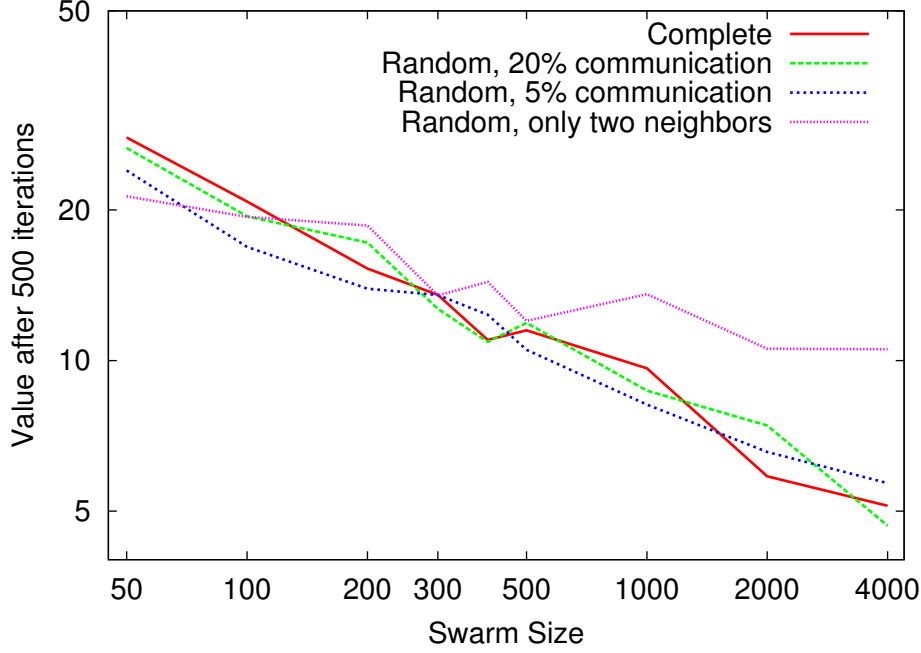


Figure 4: Random with various amounts of communication versus the Complete topology on the function Rastrigin.

With more than 4000 particles the algorithm becomes very computationally expensive on a single machine. We have implemented a parallel version of particle swarm optimization using Google’s MapReduce framework for parallel computation [3]. Soon we will have it working well enough to run swarms of hundreds of thousands of particles on hundreds of machines at a time, and we will further explore the issues involved with parallel particle swarm optimization.

Many new possibilities for topologies arise when the algorithm is run in parallel with large swarms. We have proposed, but not yet fully tested, a topology of subswarms, where each of a hundred machines would evaluate 500 or 1000 particles and send only its best value to all other machines, drastically reducing the amount of inter-processor communication. In our formal notation, the subswarms topology can be described as

$$\text{Subswarms} : N_i = \{j | j \% k = i \% k\}$$

where k is the number of subswarms in the swarm, and $\%$ is the mod operator. For a small example, in a swarm of 12 particles with 3 subswarms, the sets of neighbors would be $\{1, 4, 7, 10\}$, $\{2, 5, 8, 11\}$, and $\{3, 6, 9, 12\}$.

Our preliminary results show that this approach is very promising [6]. Figure 5 shows the results of using the Subswarms topology with 400 particles in each swarm, versus just using 400 particles in a Complete topology. Subswarms clearly outperforms Complete in this case.

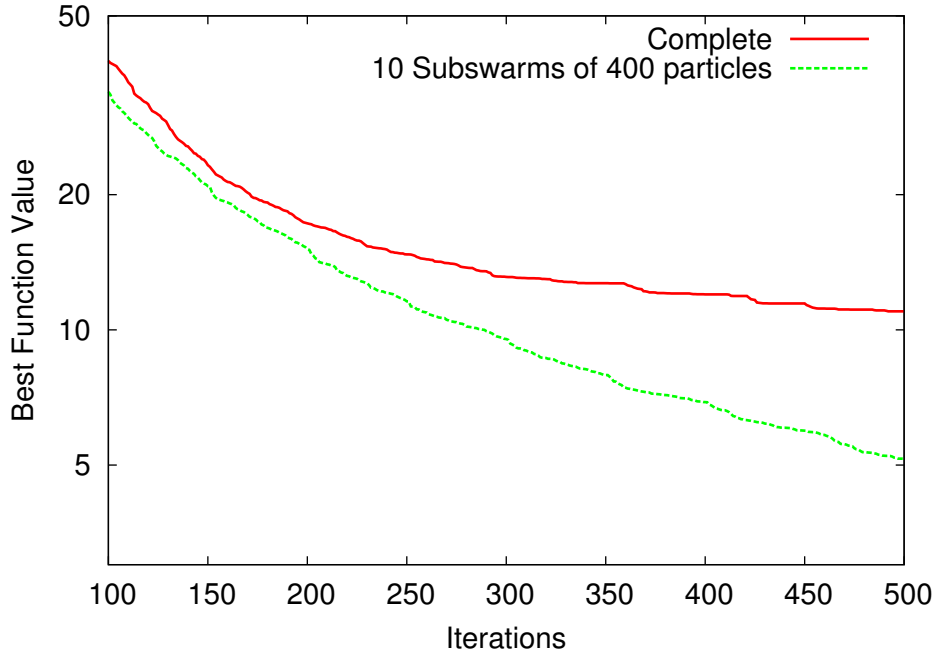


Figure 5: Subswarms versus Complete topologies on the function Rastrigin.

4 A Different Speed up—Speculative Execution

4.1 Speculative Execution in General

Speculative execution in a program is the execution of code that may or may not end up actually being needed. Modern processors routinely do this when a conditional branch is encountered—they try to predict which branch will end up being needed and continue their execution. If it turns out that the prediction was incorrect, the work is discarded. But if the prediction was right, the program can execute much faster than if it had waited on the branch.

4.2 Speculative Execution in Particle Swarm Optimization

In many optimization algorithms speculative execution is impossible, because the next iteration of the algorithm depends on the evaluation of the function during the previous iteration. To know which positions to speculatively evaluate you have to already have evaluated the current position, and then it's not speculative anymore. In simulated annealing, for instance, the probability of accepting a new position depends on the difference in function value between the new position and the current one. In gradient descent, the gradient of the function must be determined (which includes actually evaluating it) to know how far to move before sampling the function again. It may be possible to modify these algorithms to allow for speculative execution, but their standard executions cannot be done speculatively.




Figure 6: Best value found on the function Sphere versus the number of function evaluations with varying swarm sizes.

Particle swarm optimization is rather unique in that it does not require the evaluation of the function to know what the next position of each particle will be. It simply depends on *which* particle ended up having the best position and whether or not the current particle found a better spot than it had seen before. That means you can just compute all possible next positions—all combinations of which particle is the new gbest and whether or not pbest was updated (\vec{g} and \vec{p} from equation (1))—and evaluate each position in parallel. That computes two iterations at the same time.

4.3 Is it worth it?

Speculative execution uses a lot of extra work to compute two iterations at the same time, but for some functions performing more iterations makes more progress than adding extra particles. The No Free Lunch theorem again shows its head here, though, because if you do better with one algorithm on a certain function or set of functions, that's because you do worse on other functions. There are, however, a set of functions that do better with more iterations than with more particles.

Figure 6 shows an example of this: the function Sphere. The graph shows the best function value found over a number of function evaluations. This is different than previous graphs that have shown function value against a number of iterations. At each iteration, a swarm of 50 particles performs 50 function evaluations, where a swarm of 4000 particles performs 4000 evaluations. Figure 6 shows that the extra 3950 function evaluations spent

on the extra particles in a 4000 particle swarm really don't do any good. It would be more worth it to just do more iterations of a 50 particle swarm than to have a swarm of 4000 particles.

The benefit of adding more particles comes in when you have hundreds or thousands of processors in a parallel cluster. You can't use 4000 machines to run a swarm of 50 particles—3950 of the processors would be idle the whole time. Adding more particles to the swarm then improves performance without adding any extra time (ignoring inter-processor communication cost for the moment). But for some functions, a better idea than using the extra 3950 processors to add more particles to the swarm is to use the processors speculatively so that two iterations can be done at the same time. Figure 6 shows just one example of a function that would be better optimized by performing speculative execution than by adding extra particles.

5 Conclusion

In a world where people are swimming in processing power, many old algorithms need to be rethought. Parallel computing architectures allow for new techniques that were previously impossible. In the realm of optimization, particle swarm optimization is uniquely capable of reaping great benefits from parallel processing. I have reviewed just a few ways in which PSO can be improved to harness the increased power of parallel clusters. As more researchers tackle the problems of modifying and adapting algorithms to better utilize the resources that are now available, our ability to solve problems with computers will increase dramatically.

References

- [1] Daniel Bratton and James Kennedy. Defining a standard for particle swarm optimization. In *Proceedings of the IEEE Swarm Intelligence Symposium*, pages 120–127, 2007.
- [2] Maurice Clerc and James Kennedy. The particle swarm - explosion, stability, and convergence in a multidimensional complex space. *IEEE Transactions on Evolutionary Computation*, 6(1):58–73, 2002.
- [3] Jeffrey Dean and Sanjay Ghemawat. MapReduce: Simplified data processing on large clusters. *Sixth Symposium on Operating System Design and Implementation*, November 2004.
- [4] James Kennedy and Russell C. Eberhart. Particle swarm optimization. In *International Conference on Neural Networks IV*, pages 1942–1948, Piscataway, NJ, 1995.
- [5] James Kennedy and Rui Mendes. Population structure and particle swarm performance. In *Proceedings of the Congress on Evolutionary Computation (CEC 2002)*, Honolulu, Hawaii, 2002.

- [6] Andrew W. McNabb, Matthew J. Gardner, and Kevin D. Seppi. An exploration of topologies and communication in large particle swarms. In *Proceedings of the IEEE Congress on Evolutionary Computation*, May 2009.
- [7] Riccardo Poli. Dynamics and stability of the sampling distribution of particle swarm optimisers via moment analysis. *Journal of Artificial Evolution and Applications*, 8(2):1–10, 2008.
- [8] D.H. Wolpert and W.G. Macready. No free lunch theorems for optimization. *Evolutionary Computation, IEEE Transactions on*, 1(1):67–82, 1997.